

ARUNIMA SINGH THAKUR

PP LAB 6

31, C, CSE

180905218

7TH MAY 2021

Question1

Write a program in CUDA to add two Matrices for the following specifications:

- a. Each row of resultant matrix to be computed by one thread.**
- b. Each column of resultant matrix to be computed by one thread.**
- c. Each element of resultant matrix to be computed by one thread.**

```
%%cu
#include<stdio.h>
#include<cuda.h>
#include<stdlib.h>

__global__ void AddRow(int *a,int *b,int *c,int m,int n)
{
    int id=threadIdx.x;
    if(id<m)
    {
        for(int i=0;i<n;i++)
        {
            c[id*n+i]=a[id*n+i]+b[id*n+i];
        }
    }
}

__global__ void AddCol(int *a,int *b,int *c,int m,int n)
{
    int id=threadIdx.x;
    if(id<n)
    {
        for(int i=0;i<m;i++)
        {
            c[id+i*n]=a[id+i*n]+b[id+i*n];
        }
    }
}
```

```

}

__global__ void AddElement(int *a,int *b,int *c,int m,int n)
{
    int id=threadIdx.x;
    if(id<m*n)
    {
        c[id]=a[id]+b[id];
    }
}

int main()
{
    int m=3,n=3;
    //declare matrices
    int a[3][3]={ {1,1,1},{2,2,2},{3,3,3}};
    int b[3][3]={ {1,1,1},{2,2,2},{3,3,3}};
    int c[m][n];
    int *d_a,*d_b,*d_c;
    //allocate memory
    int size=m*n*sizeof(int);
    cudaMalloc((void**)&d_a,size);
    cudaMalloc((void**)&d_b,size);
    cudaMalloc((void**)&d_c,size);
    //copy from host to device
    cudaMemcpy(d_a,a,size,cudaMemcpyHostToDevice);
    cudaMemcpy(d_b,b,size,cudaMemcpyHostToDevice);

    printf("Matrix A:\n");
    for(int i=0;i<m;i++)
    {
        for(int j=0;j<n;j++)
        {
            printf("%d ",a[i][j]);
        }
        printf("\n");
    }
    printf("Matrix B:\n");
    for(int i=0;i<m;i++)
    {
        for(int j=0;j<n;j++)
        {
            printf("%d ",b[i][j]);
        }
        printf("\n");
    }
}

```

```

// Launch add() kernels on GPU
//Number of blocks=1 threads=m
printf("\n 1 a) Add row per Thread\n");
AddRow<<<1,m>>>(d_a, d_b, d_c,m,n);
cudaMemcpy(c, d_c,size,cudaMemcpyDeviceToHost);
for(int i=0;i<m;i++)
{
    for(int j=0;j<n;j++)
    {
        printf("%d ",c[i][j]);
    }
    printf("\n");
}
//Number of blocks=1 threads=n
printf("\n 1 b)Add column per Thread\n");
AddCol<<<1,n>>>(d_a, d_b, d_c,m,n);
cudaMemcpy(c, d_c,size,cudaMemcpyDeviceToHost);
for(int i=0;i<m;i++)
{
    for(int j=0;j<n;j++)
    {
        printf("%d ",c[i][j]);
    }
    printf("\n");
}
//Number of blocks=1 threads=m*n
printf("\n 1 c)Add element per Thread\n");
AddElement<<<1,m*n>>>(d_a, d_b, d_c,m,n);
cudaMemcpy(c, d_c,size,cudaMemcpyDeviceToHost);
for(int i=0;i<m;i++)
{
    for(int j=0;j<n;j++)
    {
        printf("%d ",c[i][j]);
    }
    printf("\n");
}
}

```

```

↳ Matrix A:
1 1 1
2 2 2
3 3 3
Matrix B:
1 1 1
2 2 2
3 3 3

1 a) Add row per Thread
2 2 2
4 4 4
6 6 6

1 b)Add column per Thread
2 2 2
4 4 4
6 6 6

1 c)Add element per Thread
2 2 2
4 4 4
6 6 6

```

Question2

Write a program in CUDA to multiply two Matrices for the following specifications:

- a. Each row of resultant matrix to be computed by one thread.
- b. Each column of resultant matrix to be computed by one thread
- c. Each element of resultant matrix to be computed by one thread.

```

%%cu
#include "cuda_runtime.h"
#include "device_launch_parameters.h"
#include <stdlib.h>
#include<stdio.h>
#define M1 6
#define N1 5
#define M2 5
#define N2 4
__global__ void mul_row_thread (int *A, int *B, int *C)
{
    int idx =threadIdx.x;

```

```

for(int j=0;j<N2;j++)
{
    int sum=0;
    for(int k=0;k<N1;k++)
    {
        sum+=A[idx*N1+k]*B[k*N2+j];
    }
    C[idx*N2+j] = sum;
}
}

__global__ void mul_col_thread (int *A, int *B, int *C)
{
    int idx =threadIdx.x;
    for(int j=0;j<M1;j++)
    {
        int sum=0;
        for(int k=0;k<N1;k++)
        {
            sum+=A[idx*N1+k]*B[k*N2+j];
        }
        C[idx*N2+j] = sum;
    }
}

__global__ void mul_ele_thread (int *A, int *B, int *C)
{
    int col=threadIdx.x;
    int row=blockIdx.x;
    int sum=0;
    for(int k=0;k<N1;k++)
    {
        sum+=A[row*N1+k]*B[k*N2+col];
    }
    C[row*N2+col]=sum;
}

int main ()
{
    // Host copies of the variables
    int A[M1][N1], B[M2][N2], C1[M1][N2],C2[M1][N2],C3[M1][N2];
    if(N1!=M2)
    {
        printf("Dimensions mismatch\n");
        exit(0);
    }
    for(int i=0;i<M1;i++)
    {
        for(int j=0;j<N1;j++)

```

```

{
A[i][j]=rand()%10;
}
}
for(int i=0;i<M2;i++)
{
for(int j=0;j<N2;j++)
{
B[i][j]=rand()%10;
}
}

int *d_a, *d_b, *d_c1,*d_c2 ,*d_c3;
int sizeA = sizeof(int) * M1 * N1;
int sizeB = sizeof(int) * M2 * N2;
int sizeC = sizeof(int) * M1 * N2;
cudaMalloc((void**)&d_a, sizeA);
cudaMalloc((void**)&d_b, sizeB);
cudaMalloc((void**)&d_c1, sizeC);
cudaMalloc((void**)&d_c2, sizeC);
cudaMalloc((void**)&d_c3, sizeC);
cudaMemcpy(d_a, &A, sizeA, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, &B, sizeB, cudaMemcpyHostToDevice);
mul_row_thread<<<1, M1>>>(d_a, d_b, d_c1);
cudaMemcpy(&C1, d_c1, sizeC, cudaMemcpyDeviceToHost);
mul_col_thread<<<1,N2>>>(d_a,d_b,d_c2);
cudaMemcpy(&C2, d_c2, sizeC, cudaMemcpyDeviceToHost);
mul_ele_thread<<<M1,N2>>>(d_a, d_b,d_c3);
cudaMemcpy(&C3, d_c3, sizeC, cudaMemcpyDeviceToHost);
int i,j;
printf("A:\n");
for (i = 0; i < M1; ++i) {
for (j = 0; j < N1; ++j) {
printf("%d\t", A[i][j]);
}
printf("\n");
}
printf("\n");
printf("B:\n");
for (i = 0; i < M2; ++i) {
for (j = 0; j < N2; ++j) {
printf("%d\t", B[i][j]);
}
printf("\n");
}
printf("\n");
printf("A * B:\n");
printf("\na)Each Row is computed by one thread:\n");
for (i = 0; i < M1; ++i) {

```

```

    for (j = 0; j < N2; ++j) {
        printf("%d\t", C1[i][j]);
    }
    printf("\n");
}
printf("\n");
printf("\nb)Each Column is computed by one thread:\n");
for (i = 0; i < M1; ++i) {
    for (j = 0; j < N2; ++j)
    {
        printf("%d\t", C2[i][j]);
    }
    printf("\n");
}
printf("\n");
printf("\nc)Each element is computed by one thread:\n");
for (i = 0; i < M1; ++i) {
    for (j = 0; j < N2; ++j) {
        printf("%d\t", C3[i][j]);
    }
    printf("\n");
}
printf("\n");
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c1);
cudaFree(d_c2);
cudaFree(d_c3);
getchar();
return 0;
}

```



A:

3	6	7	5	3
5	6	2	9	1
2	7	0	9	3
6	0	6	2	6
1	8	7	9	2
0	2	3	7	5

B:

9	2	2	8
9	7	3	6
1	2	9	3
1	9	4	7
8	4	5	0

A * B:

a) Each Row is computed by one thread:

117	119	122	116
118	141	87	145
114	146	76	121
110	66	104	80
113	161	135	140
68	103	86	70

b) Each Column is computed by one thread:

117	119	122	116
80	116	87	145
125	101	76	121
97	64	104	80
76	104	0	0
0	0	0	0

c) Each element is computed by one thread:

117	119	122	116
118	141	87	145
114	146	76	121
110	66	104	80
113	161	135	140
68	103	86	70