

```
!apt-get --purge remove cuda nvidia* libnvidia-*
!dpkg -l | grep cuda- | awk '{print $2}' | xargs -n1 dpkg --purge
!apt-get remove cuda-*
!apt autoremove
!apt-get update
```

```
!wget https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu
!dpkg -i cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb
!apt-key add /var/cuda-repo-9-2-local/7fa2af80.pub
!apt-get update
!apt-get install cuda-9.2
```

```
!nvcc --version
```

```
!pip install git+git://github.com/andreinechaev/nvcc4jupyter.git
```

```
%load_ext nvcc_plugin
```

**ARUNIMA SINGH THAKUR**

**SECTION C**

**ROLL NO. 31**

**180905218**

**PP LAB 8**

**2nd JUNE 2021**

### 1) 1D convolution

```
%%cu
#include<stdio.h>
#include<stdlib.h>
#include"cuda_runtime.h"
#include"device_launch_parameters.h"

__global__ void convolution_1D(int *N,int *M,int *P,int Mask_width,int Width)
{
    int i=blockIdx.x*blockDim.x+threadIdx.x;
    int p=0;
    int NStartPoint=i-(Mask_width/2); //calculating starting point for ith point in P
    for(int j=0;j<Mask_width;j++)
    {
        if(NStartPoint+j>=0 && NStartPoint+j<Width)//calculating only for non ghost elements
        {
            p+=N[NStartPoint+j]*M[j];
        }
    }
}
```

```

    P[i]=p;//save in output array
}

int main()
{
    int Width=7;//width of input and output
    int Mask_width=5;//width of mask array
    //declarations
    int N[]={100,200,300,400,500,600,700};
    int M[]={1,2,3,4,5};
    int P[Width];
    int size=Width*sizeof(int);
    int *d_N,*d_M,*d_P;
    cudaEvent_t begin,end;
    //memory allocations
    cudaMalloc((void**)&d_N,size);
    cudaMalloc((void**)&d_M,Mask_width*sizeof(int));
    cudaMalloc((void**)&d_P,size);
    //copy from host to devices
    cudaMemcpy(d_N,N,size,cudaMemcpyHostToDevice);
    cudaMemcpy(d_M,M,Mask_width*sizeof(int),cudaMemcpyHostToDevice);
    cudaEventCreate(&begin);
    cudaEventCreate(&end);
    cudaEventRecord(begin,0);
    //run kernel
    convolution_1D<<<1,Width>>>(d_N,d_M,d_P,Mask_width,Width);
    cudaEventRecord(end,0);
    cudaEventSynchronize(end);
    float elapsedTime;
    cudaEventElapsedTime(&elapsedTime,begin,end);
    //copy from device to host
    cudaMemcpy(P,d_P,size,cudaMemcpyDeviceToHost);

    for(int i=0;i<Width;i++)
    {
        printf("%d\t",P[i]);
    }
    printf("\nTime taken by kernel= %f",elapsedTime);
    //free memory
    cudaFree(d_N);
    cudaFree(d_M);
    cudaFree(d_P);
    return 0;
}

```

## OUTPUT:

2600 4000 5500 7000 8500 6000 3800

Time taken by kernel= 0.028672

## 2) 1D convolution with constant memory for mask array(kernel1)and shared memory for N (kernel2)

```

%%cu
#include<stdio.h>
#include<stdlib.h>
#include"cuda_runtime.h"
#include"device_launch_parameters.h"
#define MAX_MASK_WIDTH 10

__constant__ float M[MAX_MASK_WIDTH]; //declarations of mask array in constant memory

__global__ void convo1Dconstmem(float *N,float *P,int Mask_width,int Width)
{
    int i=blockIdx.x*blockDim.x+threadIdx.x;
    int p=0;
    int NStartPoint=i-(Mask_width/2); //calculating starting point for ith point in P
    for(int j=0;j<Mask_width;j++)
    {
        if(NStartPoint+j>=0 && NStartPoint+j<Width) //calculating only for non ghost elements
        {
            p+=N[NStartPoint+j]*M[j];
        }
    }
    P[i]=p; //save in output array
}

__global__ void convo1Dsharedmem(float *N,float *P,int Mask_width,int Width)
{
    extern __shared__ float temp_array[]; //declaring a array in shared memory
    int id=blockIdx.x*blockDim.x+threadIdx.x;
    temp_array[threadIdx.x]=N[id];
    __syncthreads();
    float p=0;
    for(int i=0;i<Mask_width;i++)
    {
        if(threadIdx.x+i>=blockDim.x)
        {
            p+=N[id+i]*M[i];
        }
        else
        {
            p+=temp_array[threadIdx.x+i]*M[i];
        }
    }
    P[id]=p; //save in output array
}

int main()
{
    //declarations
    int Width=7;
    int Mask_width=5;
    float h_N[]={100,200,300,400,500,600,700};
    float h_M[]={1,2,3,4,5};
    float h_P[Width];
    int size=Width*sizeof(float);

```

```

float *d_N,*d_P;
float eTime1,eTime2;
cudaEvent_t begin,end,begin1,end1;
//memory allocations
cudaMalloc((void**)&d_N,size);
cudaMalloc((void**)&d_P,size);
//copy from host to device
cudaMemcpy(d_N,h_N,size,cudaMemcpyHostToDevice);
cudaMemcpyToSymbol(M,h_M,Mask_width*sizeof(float));
cudaEventCreate(&begin);
cudaEventCreate(&end);
cudaEventCreate(&begin1);
cudaEventCreate(&end1);
cudaEventRecord(begin,0);
//run kernel
convolDconstmem<<<1,Width>>>(d_N,d_P,Mask_width,Width);
cudaEventRecord(end,0);
cudaEventSynchronize(end);
cudaEventElapsedTime(&eTime1,begin,end);
//copy from device to host
cudaMemcpy(h_P,d_P,size,cudaMemcpyDeviceToHost);
for(int i=0;i<Width;i++)
{
    printf("%f\t",h_P[i]);
}
printf("\nTime taken by kernel 1= %f\n",eTime1);
cudaEventRecord(begin1,0);
//run kernel
convolDsharedmem<<<1,Width>>>(d_N,d_P,Mask_width,Width);
cudaEventRecord(end1,0);
cudaEventSynchronize(end1);
cudaEventElapsedTime(&eTime2,begin1,end1);
//copy from device to host
cudaMemcpy(h_P,d_P,size,cudaMemcpyDeviceToHost);
for(int i=0;i<Width;i++)
{
    printf("%f\t",h_P[i]);
}
printf("\nTime taken by kernel 2= %f",eTime2);
//free memory
cudaFree(d_N);
cudaFree(d_P);
return 0;
}

```

## OUTPUT:

```

2600.000000 4000.000000 5500.000000 7000.000000 8500.000000 6000.000000 3800.000000
Time taken by kernel 1= 0.022176 2600.000000 4000.000000 5500.000000 7000.000000
8500.000000 6000.000000 3800.000000
Time taken by kernel 2= 0.000000

```

### 3) Sparse Matrix vector multiplication (CSR format)

```

%%cu
#include<stdio.h>
#include<stdlib.h>
#include"cuda_runtime.h"
#include"device_launch_parameters.h"

__global__ void SpMV_CSR(int num_rows,int *data,int *col_index,int *row_ptr,int *x,int *y)
{
    int row=threadIdx.x;
    if(row<num_rows)
    {
        int dot=0;
        int row_start=row_ptr[row];
        int row_end=row_ptr[row+1];
        for(int i=row_start;i<row_end;i++)
        {
            dot+= data[i]*x[col_index[i]];
        }
        y[row]=dot;
    }
}

int main()
{
    //declarations
    int n=4;
    int y[n],row_ptr[n+1];
    int ipmat[n][n]={1,2,3,4},{5,6,7,8},{9,10,11,12},{13,14,15,16}};
    int x[]={100,200,300,400};
    int nonzerocount=0;
    //finding number of non zero elements and row ptr array
    for(int i=0;i<n;i++)
    {
        row_ptr[i]=nonzerocount;
        for(int j=0;j<n;j++)
        {
            if(ipmat[i][j]!=0)
            {
                nonzerocount++;
            }
            printf("%d\t",ipmat[i][j]);
        }
        printf("\n");
    }
    row_ptr[n]=nonzerocount;
    int data[nonzerocount],col_index[nonzerocount];
    int k=0;
    //finding data and col_index array
    for(int i=0;i<n;i++)
    {
        for(int j=0;j<n;j++)
        {
            if(ipmat[i][j]!=0)

```

```

        {
            data[k]=ipmat[i][j];
            col_index[k++]=j;
        }
    }
}
printf("\ndata array\t");
for(int i=0;i<nonzerocount;i++)
{
    printf("%d\t",data[i]);
}
printf("\ncol_index array\t");
for(int i=0;i<nonzerocount;i++)
{
    printf("%d\t",col_index[i]);
}
printf("\nrow_ptr array\t");
for(int i=0;i<=n;i++)
{
    printf("%d\t",row_ptr[i]);
}
printf("\nvector X\t");
for(int i=0;i<n;i++)
{
    printf("%d\t",x[i]);
}
int *d_data,*d_col_index,*d_row_ptr,*d_x,*d_y;
//memory allocations
cudaMalloc((void**)&d_data,nonzerocount*sizeof(int));
cudaMalloc((void**)&d_col_index,nonzerocount*sizeof(int));
cudaMalloc((void**)&d_row_ptr,(n+1)*sizeof(int));
cudaMalloc((void**)&d_x,n*sizeof(int));
cudaMalloc((void**)&d_y,n*sizeof(int));
//copy from host to device
cudaMemcpy(d_data,data,nonzerocount*sizeof(int),cudaMemcpyHostToDevice);
cudaMemcpy(d_col_index,col_index,nonzerocount*sizeof(int),cudaMemcpyHostToDevice);
cudaMemcpy(d_row_ptr,row_ptr,(n+1)*sizeof(int),cudaMemcpyHostToDevice);
cudaMemcpy(d_x,x,n*sizeof(int),cudaMemcpyHostToDevice);
//run kernel
SpMV_CSR<<<1,n>>>(n,d_data,d_col_index,d_row_ptr,d_x,d_y);
//copy from device to host
cudaMemcpy(y,d_y,n*sizeof(int),cudaMemcpyDeviceToHost);

printf("\nresult\t");
for(int i=0;i<n;i++)
{
    printf("%d\t",y[i]);
}
//free memory
cudaFree(d_data);
cudaFree(d_col_index);
cudaFree(d_row_ptr);
cudaFree(d_x);
cudaFree(d_y);
return 0;

```

```
}
```

## OUTPUT:

```
1 2 3 4
```

```
5 6 7 8
```

```
9 10 11 12
```

```
13 14 15 16
```

```
data array 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
```

```
col_index array 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3
```

```
row_ptr array 0 4 8 12 16
```

```
vector X 100 200 300 400
```

```
result 3000 7000 11000 15000
```

## 4) Matrix multiplication using 2D grids and 2D blocks

```
%%cu
#include<stdio.h>
#include<stdlib.h>

__global__ void matrixMul(const int *a, const int *b, int *c, int m,int n,int o)
{
    //row and col calculations
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    c[row * o + col] = 0;
    //calculating one element
    for (int k = 0; k < n; k++) {
        c[row * o + col] += a[row * n + k] * b[k * o + col];
    }
}

int main()
{
    //declarations
    int size =sizeof(int);
    int m=4,n=2,o=4;
    int a[m][n];
    int b[n][o];
    int c[m][o];
    for(int i=0;i<m;i++)
    {
        for(int j=0;j<n;j++)
        {
            a[i][j]=1+i+j;
        }
    }
    for(int i=0;i<n;i++)
    {
        for(int j=0;j<o;j++)
```

```

    {
        b[i][j]=2+i+j;
    }
}
int *d_a, *d_b, *d_c;
//memory allocations
cudaMalloc(&d_a,m*n*size);
cudaMalloc(&d_b,n*o*size);
cudaMalloc(&d_c,m*o*size);
//copy from host to device
cudaMemcpy(d_a,a,m*n*size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b,b,n*o*size, cudaMemcpyHostToDevice);
//dimensions for grid and block
int thread=2;
dim3 threads(thread,thread);
dim3 blocks((m*o)/(4*thread),(m*o)/(4*thread));

//run kernel
matrixMul<<<blocks, threads>>>(d_a, d_b, d_c,m,n,o);

// copy from device to host
cudaMemcpy(c, d_c,m*o*size, cudaMemcpyDeviceToHost);
printf("Matrix A\n");
for(int i=0;i<m;i++)
{
    for(int j=0;j<n;j++)
    {
        printf("%d\t",a[i][j]);
    }
    printf("\n");
}
printf("Matrix B\n");
for(int i=0;i<n;i++)
{
    for(int j=0;j<o;j++)
    {
        printf("%d\t",b[i][j]);
    }
    printf("\n");
}
printf("Matrix C\n");
for(int i=0;i<m;i++)
{
    for(int j=0;j<o;j++)
    {
        printf("%d\t",c[i][j]);
    }
    printf("\n");
}
//free memory
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

return 0;

```



```
}
```

**OUTPUT:**

Matrix A 1 2

2 3

3 4

4 5

Matrix B 2 3 4 5

3 4 5 6

Matrix C 8 11 14 17

13 18 23 28

18 25 32 39

23 32 41 50

---

✓ 0s completed at 1:48 PM

